
AMDOX TECHNOLOGIES

Enterprise AI-Powered Cloud ERP Suite

Next-Generation Intelligent Resource Planning Platform

Project Report

Industry Internship Project • Engineering Division

Project Code	AMX-ERP-2026-04
Prepared by	Pawan Nannaware
Organization	Amdox Technologies
Domain	Enterprise Software / Cloud Infrastructure
Report date	July 2026
Version	1.0

Live demo	https://erp.92-4-86-3.sslip.io
API documentation	erp.92-4-86-3.sslip.io/api-docs
Demo video	LinkedIn — full-stack development walkthrough
Identity provider	kc.92-4-86-3.sslip.io (Keycloak)

Contents

1	Executive Summary	3
2	Problem Statement & Objectives	3
2.1	Target Users	3
2.2	Success Criteria	4
3	System Architecture	4
3.1	Topology	4
3.2	Monorepo	4
3.3	Multi-Tenant Design	5
3.4	Design Patterns in Use	5
3.5	Data Model	5
4	Functional Modules	5
4.1	Authentication & Identity (F-01)	5
4.2	Financial Ledger — GL (F-02)	8
4.3	AP / AR Automation (F-03)	9
4.4	HR & Payroll (F-04)	10
4.5	Supply Chain Management (F-05)	12
4.6	AI Demand Forecasting (F-06)	12
4.7	Project Management (F-07)	13
4.8	Business Intelligence (F-08)	13
4.9	Audit & Compliance (F-09)	14
4.10	Notification Engine (F-10)	14
4.11	API Gateway & Search (F-11)	14
5	Cross-Module Integration	14
5.1	Procure-to-Pay (P2P)	14
5.2	Hire-to-Pay (HR → Finance) & Order-to-Cash (Sales → Finance)	15
5.3	Background Job Architecture	15
6	Security & Compliance	15
6.1	Authentication & Token Lifecycle	15
6.2	Role-Based Access Control	16
6.3	Transport & Headers	16
6.4	Rate Limiting & Throttling	16
6.5	Static Analysis & Supply-Chain Scans	16
6.6	Key Vulnerabilities Found & Fixed	17
7	Quality Assurance & Performance	17
7.1	API Testing — 64/64 Authenticated Checks	17
7.2	Unit Tests — 31 CI-Enforced Tests	17
7.3	Load Testing — k6 at 2,000 VUs	17
7.4	Lighthouse Audits	18
7.5	Database Query Optimisation	18
8	Deployment & DevOps	19
8.1	Production Environment	19
8.2	CI Pipeline	19
8.3	Container & Kubernetes	19
9	Observability	19

10 Engineering Challenges & Notable Fixes 21

11 Requirements Compliance & Known Deviations 21

11.1 Functional Requirements Scoreboard 22

11.2 Non-Functional Requirements — Evidence 22

11.3 Technology Stack — Spec vs. Delivered 23

11.4 Known Deviations 23

12 Learnings & Conclusion 23

12.1 Key Technical Learnings 23

12.2 Roadmap 24

12.3 Conclusion 24

1 Executive Summary

This report documents the design, implementation, verification, and deployment of the **Amdox Enterprise AI-Powered Cloud ERP Suite** — a scalable, multi-tenant SaaS platform delivering financial management, supply chain automation, HR & payroll, project tracking, and self-serve business intelligence, augmented with machine-learning demand forecasting. The system was built as a Turborepo monorepo over a 28-day execution plan defined in the company’s technical design document (TDD, AMX-ERP-2026-04).

Outcome: 11 of the 12 functional requirements (F-01–F-12) are implemented, verified against the running system, and publicly deployed at <https://erp.92-4-86-3.sslip.io> with automatic HTTPS. The single de-scoped requirement (F-12 Offline/PWA) is tracked on the roadmap, following the TDD’s own Day-1 mandate to “define MVP feature scope and de-scope list.” Of the 43 tracked engineering tasks, 39 are complete (Frontend 18/18, Integrations 9/9, Backend 9/11, Platform 3/5).

The platform is not a collection of isolated CRUD screens: departments coordinate through domain events, so a project manager’s material request becomes a purchase requisition, an approved purchase order notifies the vendor by e-mail and webhook, a goods receipt creates FIFO cost layers and triggers a three-way-matched AP invoice, and the general ledger and project budgets update themselves — with every mutation recorded in a tamper-evident, hash-chained audit log.

Project at a Glance	
Scope delivered	11/12 functional requirements; 6 business modules; 43 application screens
Data model	67 Prisma models on PostgreSQL 17, tenant-scoped by an auto-injecting ORM client
Verification	64/64 authenticated API checks; 31 CI-enforced unit tests on the money paths; k6 load test at 2,000 concurrent VUs; Lighthouse ≥ 90 on audited pages
Security	Realm-per-tenant SSO with per-tenant MFA; RBAC on all business controllers; CI static-analysis gate for tenant isolation; TruffleHog/Grype/Trivy scans
Operations	Live public deployment (Oracle Cloud, Caddy TLS, pm2); Docker/Helm/ArgoCD GitOps with Istio canary; full OpenTelemetry + Prometheus/Grafana/Loki/Tempo observability with SLA alerting

2 Problem Statement & Objectives

Mid-market and enterprise organisations typically run finance, HR, procurement, and project delivery on separate legacy systems. The TDD frames the consequence directly: data silos across “6+ legacy systems,” manual approval cycles measured in days, and reporting that requires IT tickets rather than self-service. The objective set for this project was to deliver a cloud-native ERP suite that consolidates these fragmented workflows into a single intelligent platform — reducing operational overhead, enabling data-driven decisions, and accelerating time-to-insight for business leaders.

2.1 Target Users

The platform serves six user segments, each with a dedicated, role-gated workspace:

User Segment	Primary Use Case (per TDD)
C-Suite / Executives	Real-time dashboards, KPI monitoring, board-level reporting
Finance Teams	GL management, AP/AR automation, multi-currency reconciliation
HR & Payroll Teams	Employee lifecycle, attendance, payroll processing, compliance
Supply Chain Managers	Procurement, inventory, vendor management, demand planning
Project Managers	Resource allocation, milestone tracking, budget management
IT Administrators	Tenant configuration, SSO, audit logs, security policies

2.2 Success Criteria

Success was defined by the TDD’s acceptance criteria per functional requirement (F-01–F-12), its non-functional targets (99.9% availability, <300 ms P95 API latency, 2,000 concurrent users per tenant, OWASP Top 10 alignment), and five submission deliverables: this report, a live public demo URL, the GitHub repository, a comprehensive README, and a scenario-based demo video. Section 11 maps the delivered system against every one of these criteria, including honest deviations.

3 System Architecture

3.1 Topology

The platform is a **modular monolith with one satellite microservice**: a NestJS 11 API hosting seventeen domain modules behind a single versioned gateway (`api/v1`), plus a Python FastAPI service dedicated to machine learning. This follows the TDD’s “in-process modules (monolith) → future microservices” guidance — module boundaries are enforced by events rather than by network hops, keeping cross-module workflows transactional where it matters (Section 5).




Figure 1: Deployed system topology. Everything runs on one Oracle Cloud VM: web/API/ml-service as pm2-managed production builds, the data and identity layer as Docker containers, and the observability stack as a separate Docker Compose project.

3.2 Monorepo

The repository is a Turborepo + pnpm workspace: `apps/web` (43 screens in Next.js route groups), `apps/api`, `apps/ml-service`, and `packages/db` — the shared Prisma schema and the tenant-aware

database client every module consumes. `infra/` holds the operational estate: Docker Compose stacks (dev, prod, observability), the Helm chart, ArgoCD manifests, and pm2 process configs. Conventional commits are enforced by commitlint + Husky, and ESLint/Prettier run repo-wide with zero violations.

3.3 Multi-Tenant Design

Tenancy is enforced at three independent layers, so a single mistake cannot leak data across companies:

- 1. **Identity** — every tenant gets its *own Keycloak realm*, created programmatically at signup, with realm roles (SuperAdmin, TenantAdmin, Manager, Viewer, Employee). Login discovers the home realm from the tenant slug; Google/SSO logins auto-link to existing accounts by verified e-mail, and a per-tenant **MFA toggle** forces OTP enrolment for every user in that tenant.
- 2. **Request context** — a `TenantContextInterceptor` extracts the tenant from the validated JWT and stores it in `AsyncLocalStorage`, making it available to any code on the request path without parameter threading.
- 3. **Data** — the shared Prisma client is wrapped with `$extends()` logic that *auto-injects* `tenantId` into every query’s `where/data`. All 38 services use this client, and a custom TypeScript-AST static analyzer runs **in CI as a permanent gate**, failing the build if any query lacks tenant scoping (Section 6 covers the two real cross-tenant bugs it caught).

Beyond isolation, tenants are commercially separable: a `ModuleGuard` reads per-tenant module licensing, so individual ERP modules can be switched on or off per customer — an extra the TDD did not require but which makes the platform genuinely sellable as SaaS.

3.4 Design Patterns in Use

Pattern	Where it earns its keep
Domain events	Modules never call each other directly: SCM emits <code>goods.received</code> ; Finance’s bridge listener reacts. Segregation of duties falls out of the design — SCM cannot approve or pay invoices.
Outbox	Finance writes <code>OutboxEvent</code> rows in the <i>same transaction</i> as the business change; a <code>BullMQ</code> worker delivers them — no lost events on a crash.
Saga w/ compensation	Payroll runs process 500-employee chunks asynchronously; a failed run <i>wipes its partial payslips</i> before retrying (the fix proven by <code>verify:payroll-retry</code>).
Guards & decorators	<code>@Roles()</code> + <code>RolesGuard</code> on all business controllers (78 declarations); <code>ModuleGuard</code> for licensing; global validation pipe with class-validator DTOs.
Auto-posting listeners	GL subscribes to <code>invoice.approved</code> , <code>payment.made</code> , <code>payroll.completed</code> , ... and posts balanced journal entries idempotently (duplicate-checked on source module + ID).

3.5 Data Model

The schema comprises **67 Prisma models** on PostgreSQL 17, organised into ten domains (Figure 2). Conventions applied throughout: UUID keys, `Decimal(18,4)` for all money, UTC timestamps, soft deletes, and composite indexes added where an 80,000-row `EXPLAIN ANALYZE` audit proved they pay off.

4 Functional Modules

4.1 Authentication & Identity (F-01)

The auth module integrates **Keycloak 25** as the identity broker. Every new tenant triggers a call to the Keycloak Admin API that creates a dedicated realm, provisions the five system roles




Figure 2: Database model — 67 models, 77 relations, grouped by domain. Every model carries `tenantId`; tenant-scoping edges are omitted for clarity.

(SuperAdmin, TenantAdmin, Manager, Viewer, Employee), and configures the OIDC client for the web app. Three login paths converge into a single token:

- **Path A — Username/password:** standard Keycloak login form.
- **Path B — Personal Google:** Google OAuth identity provider added per realm; on first login an account is created and auto-linked by verified email so the user never sees a conflict screen.
- **Path C — Corporate SSO:** Keycloak's broker supports SAML and OIDC external IdPs (Azure AD, Google Workspace) — plumbing is in place and tested with a Keycloak-to-Keycloak federation.

After login, every API call carries a short-lived RS256 JWT. The `KeycloakStrategy` validates the signature against the realm's JWKS endpoint, maps the `ssoSubject` to a local `User` record, and loads tenant + roles into `req.user`. A **per-tenant MFA toggle** (Settings → Authentication) adds `Configure OTP` as a required action for every user in that realm, enforcing TOTP on both password and SSO login paths. Token blacklisting on logout is implemented via a Redis set that every guarded route checks.




Figure 3: Authentication flow. Paths A/B/C are three ways to prove identity; everything from the merge point on — MFA, tokens, per-request guards — is identical regardless of path (adapted from `docs/auth-flow.md`).

4.2 Financial Ledger — GL (F-02)




Figure 4: GL journal-entry creation screen. Debit/credit balancing is enforced by both the UI and the `GLService` before persistence.

The General Ledger (`gl.service.ts`) enforces three accounting invariants at the service layer:

1. **Double-entry:** `createJournalEntry` rejects any entry where $\Sigma \text{Debits} \neq \Sigma \text{Credits}$ with a 400 error before touching the database.
2. **Period-close guard:** posting into a locked `FiscalPeriod` throws; periods auto-open on the first business day of each month via cron.
3. **Idempotency:** each auto-posting listener checks (`sourceModule`, `sourceId`) before inserting, so a retry never creates a duplicate journal entry.

Auto-posting listeners handle every financial event:

Event	Debit	Credit
invoice.approved (AP)	1300 Inventory	2000 Accounts Payable
invoice.issued (AR)	1200 Accounts Receivable	4000 Revenue
payment.received	1000 Cash	1200 Accounts Receivable
payment.made	2000 Accounts Payable	1000 Cash
payroll.completed	6000 Salary Expense	2100 Payroll Payable

Table 1: Automatic journal entries posted by the GL service on domain events.

Foreign-exchange rates are fetched daily from the ECB feed (29 currencies, per-tenant `ExchangeRate` rows); actual conversion in postings is on the roadmap.

4.3 AP / AR Automation (F-03)

INVOICE #	VENDOR	PO	ISSUED	DUE DATE	AMOUNT	STATUS	ACTION
INV-OCR-4621	—	—	10/7/2026	9/8/2026	₹1,500	PENDING MATCH	Approve Cancel
INV-OCR-9412	—	—	10/7/2026	9/8/2026	₹1,500	CANCELLED	
INV-OCR-9964	—	—	10/7/2026	9/8/2026	₹1,500	CANCELLED	
INV-OCR-2134	—	—	10/7/2026	9/8/2026	₹1,500	CANCELLED	
INV-AUTO-17837830545754	Chai Break Pantry Supplies	PO-1783783285512	10/7/2026	9/8/2026	₹8,000	APPROVED	
INV-AUTO-1783676935393	DailyBrew Office Refreshments	PO-1783675912659	10/7/2026	9/8/2026	₹9,000	PAID	
INV-AUTO-1783676934337	Fresh Bites Corporate Pantry	PO-1783675914184	10/7/2026	9/8/2026	₹9,600	APPROVED	
INV-AUTO-1783676933824	Chai Break Pantry Supplies	PO-1783675915868	10/7/2026	9/8/2026	₹6,400	APPROVED	
INV-AUTO-1783676929918	Chai Break Pantry Supplies	PO-1783675915814	10/7/2026	9/8/2026	₹5,600	APPROVED	
INV-AUTO-1783674517948	Office Supplies Co	PO-1783674478899	10/7/2026	9/8/2026	₹3,00,000	APPROVED	

Figure 5: Accounts Payable invoice list. Statuses (PENDING_MATCH, APPROVED, PAID) reflect 3-way match outcomes; the upload button triggers AWS Textract OCR.

Accounts Payable

Invoice processing supports two entry paths:

- **Auto-creation:** when a goods receipt is committed, the SCM–Finance bridge creates an AP invoice for the received quantities and immediately runs the 3-way match.
- **Manual upload:** finance staff uploads a PDF; AWS Textract `AnalyzeExpense` extracts vendor, date, line items, and totals with per-field confidence scores (a local mock is the fallback when no AWS credentials are configured).

The **three-way match** (`performThreeWayMatch`) verifies: (i) invoice vendor matches PO vendor; (ii) per-line invoiced quantity $\leq$ received quantity; (iii) per-line unit price within $\pm 2\%$ of the PO line price (header-total fallback when lines cannot be resolved to products). A passing match auto-approves the invoice in under one second; a failure parks it as `PENDING_MATCH` for manual review. Payment runs batch-settle approved invoices with per-invoice failure isolation and `PARTIALLY_PAID` states.




Figure 6: Three-way match decision flow (performThreeWayMatch). Line-level checking is preferred; header-total is a fallback when invoice lines cannot be mapped to PO products.

Accounts Receivable

Customers, AR invoices, payment recording, and an **aging report** (grouping outstanding balances by 0–30, 31–60, 61–90, and >90 days) are all implemented. A sales-order module provides the order-to-cash entry point: sales order → AR invoice → GL posts revenue.

4.4 HR & Payroll (F-04)

The screenshot shows the 'Payroll' dashboard for July 2026. It displays summary statistics and a detailed table of employee payroll data.

EMPLOYEE	PAY PERIOD	GROSS PAY	DEDUCTIONS	NET PAY	STATUS	PAYSLIP
MD Manish Desai	Jul 2026	₹75,000	-₹15,820	₹59,180	Processed	↓ Payslip
AM Arjun Malhotra	Jul 2026	₹1,20,000	-₹35,740	₹84,260	Processed	↓ Payslip
VC Varun Chopra	Jul 2026	₹55,000	-₹11,660	₹43,340	Processed	↓ Payslip
RD Rohan Das	Jul 2026	₹1,20,000	-₹35,740	₹84,260	Processed	↓ Payslip
VM Vikram Mehta	Jul 2026	₹1,20,000	-₹35,740	₹84,260	Processed	↓ Payslip
KR Kavya Reddy	Jul 2026	₹95,000	-₹19,980	₹75,020	Processed	↓ Payslip
PS Pooja Saxena	Jul 2026	₹75,000	-₹15,820	₹59,180	Processed	↓ Payslip
ST Sara Thomas	Jul 2026	₹48,000	-₹5,980	₹42,020	Processed	↓ Payslip

Summary statistics from the dashboard:

- TOTAL GROSS: ₹30,61,65,354
- NET PAY: ₹21,58,30,149.216
- DEDUCTIONS: ₹9,03,35,204.784
- PROCESSED: 69 / 69

Figure 7: Payroll run dashboard. Each row shows gross pay, statutory deductions, and net pay; the PDF payslip is downloadable from the same screen.

Employee Management

Full CRUD with employment contracts (salary history), department hierarchy, `managerId` tree, and soft-delete. Employees self-serve via `GET /hr/employees/me` and a personal home-page dashboard (tabbed: Profile / Attendance & Leave / Pay).

Leave Management

A leave state-machine enforces: only the *direct manager*, an *HR-department approver*, or a *TenantAdmin* can approve or reject a leave request. The state-machine is pure-function, making approval rules easy to unit-test in isolation.




Figure 8: Leave request state machine (`LeaveStateMachine.validateTransition`). Only PENDING is a valid starting state; APPROVED and REJECTED are terminal. Every transition re-checks authorization — the caller must be the employee’s direct manager, an HR-department approver, or a TenantAdmin, or the service throws `ForbiddenException`. CANCELLED is defined in the schema but has no reachable code path — a documented dead state, not a missing feature.

Attendance

Clock-in/out records accumulate per-day overtime minutes; the payroll engine converts these into overtime pay at $1.5\times$ the hourly rate.

Payroll Engine

A BullMQ worker processes runs asynchronously in **chunks of 500 employees** for memory safety. The gross-to-net computation applies: PF (12% employee + 12% employer on basic), ESI (0.75% employee + 3.25% employer, under the Rs. 21,000 ceiling), Professional Tax, Labour Welfare Fund, gratuity accrual, and TDS from per-tenant configurable tax slabs. Generated payslip PDFs are stored in MinIO (S3-compatible). On completion, the `payroll.completed` event posts the salary GL entry and credits each linked project’s labour-cost budget.




Figure 9: Payroll saga compensation. Before the fix, a mid-run crash left chunk 0–2’s payslips in place and the retry re-inserted them — 21 employees produced 42 payslips. The compensation step (delete-before-reprocess) makes the retry idempotent, verified live with `pnpm -filter api verify:payroll-retry`.

4.5 Supply Chain Management (F-05)




Figure 10: Purchase Order list with status badges (DRAFT, APPROVED, RECEIVED). The Goods Receipt button opens the per-line quantity capture dialog.

The SCM module spans the full procure-to-pay cycle:

- **Vendors & Products:** full CRUD; a vendor portal (pre-shared portal-key header auth, scoped to that vendor's own data) lets suppliers view their POs and post acknowledgements without an internal user account.
- **Requisition → PO:** purchase requisitions (from project managers or the reorder engine) are reviewed by supply-chain managers, who convert them to POs against a chosen vendor.
- **Goods Receipt:** supports full and *partial* delivery — each receipt records per-line quantities (GoodsReceiptLine), enabling the PO to remain **PARTIALLY_RECEIVED** across multiple deliveries.
- **FIFO Inventory:** each goods receipt creates a cost layer; outbound stock movements consume oldest layers first, giving accurate weighted-average cost of goods sold.
- **Reorder Automation:** configurable reorder points per product/warehouse; the scheduler auto-drafts a purchase requisition whenever stock falls below the threshold.

4.6 AI Demand Forecasting (F-06)

The **ml-service** (Python FastAPI) runs two models per SKU: **Prophet** for products with seasonal/trend patterns, and a **PyTorch LSTM** for high-volume SKUs. Predictions are cached in Redis and served to the API via internal REST. A weekly **forecast-retrain** BullMQ job pulls fresh sales and stock data, retrains models, and invalidates the cache. If MAPE exceeds 12%, the service emits a **forecast.mape_breach** event that triggers an alert to supply-chain managers.




Figure 11: Demand forecast chart for a sample SKU — 90-day prediction with confidence bands (80% and 95%), actual vs. predicted overlay, and model MAPE display.

4.7 Project Management (F-07)

Projects move through a lifecycle (Planning → Active → On Hold → Closed). Key capabilities include: DAG-validated task dependencies (cycle detection on save), an interactive **Gantt chart**, a resource-utilisation heatmap, and milestone tracking. Budget actuals are fed automatically from two sources: approved AP invoices whose PO carries a `projectId`, and payroll runs that credit labour costs from resource allocations. When actuals exceed budget by >10%, an overrun alert fires through the notification engine.

4.8 Business Intelligence (F-08)




Figure 12: BI drag-and-drop dashboard builder. Widgets snap to a 12-column grid; the right panel configures data source, chart type, and refresh interval.

The BI module provides a **drag-and-drop dashboard builder** (`react-grid-layout`) with 10+ chart types (line, bar, pie, scatter, KPI card, table, map). Dashboards are personal or shared at the tenant level. Widgets refresh via **Server-Sent Events** (SSE) on a configurable interval, keeping

live metrics current without polling. Scheduled reports generate real PDF/XLSX files server-side and deliver them by email on a cron schedule.

4.9 Audit & Compliance (F-09)

Every mutation in the system writes to an **audit log** protected by a SHA-256 hash chain: each row includes `previousHash` so any tampering breaks the chain. A `GET /audit/verify` endpoint re-computes hashes server-side and reports any gap. GDPR endpoints allow data-subject requests (DSR) and consent management by IT admins.

4.10 Notification Engine (F-10)

A unified notification service fans events out across four channels: **in-app** (Server-Sent Events, shown as a bell counter in the nav), **email** (Nodemailer → Mailpit in dev, SMTP in prod), **SMS**, and **HMAC-signed webhooks**. Each user configures per-channel preferences. The `notification-dispatch` BullMQ queue retries failed deliveries up to five times with exponential back-off; timed-out jobs are visible in Bull Board at `/admin/queues`.

4.11 API Gateway & Search (F-11)

- **REST API:** all endpoints under `/api/v1`, documented in Swagger UI at `/api-docs`, with request/response DTOs and role annotations.
- **GraphQL:** Apollo Server with the same Keycloak auth guard; supports flexible BI-style queries for the dashboard builder.
- **Elasticsearch:** 10 entity types are indexed in real time (employees, POs, invoices, customers, projects, leave requests, audit logs, journal entries, vendors, products). The global search bar (`Ctrl+K / /`) queries all indices simultaneously with fuzzy matching and groups results by entity. Graceful degradation to a Postgres `ILIKE` fallback when ES is unavailable.
- **Outbound webhooks:** HMAC-SHA256 signed, retried, and delivered through the notification queue.

5 Cross-Module Integration

5.1 Procure-to-Pay (P2P)

This is the flagship cross-module workflow, connecting five modules through domain events:




Figure 13: Procure-to-Pay flow. Dashed-boundary modules (Finance) cannot be directly called by SCM — communication is exclusively through domain events.

1. A **Project Manager** raises a material need from within the PM module (or the reorder engine raises it automatically when stock hits the threshold).
2. **SCM** converts the requisition to a **Purchase Order** against a chosen vendor. The vendor is notified by email and webhook; the vendor portal allows online acknowledgement.
3. **Goods Receipt:** warehouse staff record what actually arrived — full or per-line partial. `PurchaseService` commits the `GoodsReceipt`, `GoodsReceiptLines`, stock upsert, and FIFO cost layer *in one DB transaction*, then emits `goods.received`.

4. **ScmFinanceBridgeListener** (Finance) consumes the event and calls `createInvoiceFromGoodsReceipt`, which creates an AP invoice for the received quantities and immediately runs `performThreeWayMatch`.
5. A passing match **auto-approves** the invoice (<1 s); failure leaves it as `PENDING_MATCH` for manual Finance review.
6. `invoice.approved` fires, and the **GL** auto-posts Dr Inventory 1300 / Cr AP 2000. The project budget actuals increment by the invoice amount.

5.2 Hire-to-Pay (HR → Finance) & Order-to-Cash (Sales → Finance)

The other two company-wide loops are shorter chains than P2P but follow the same principle: a business event in one module ends in an automatic, balanced GL posting.




Figure 14: Hire-to-Pay (top) and Order-to-Cash (bottom) loops. Both terminate in an automatic, idempotent GL posting triggered by a domain event, the same pattern P2P uses (Figure 13).

5.3 Background Job Architecture

Queue	Producer	Consumer	Purpose
payroll	PayrollService	PayrollProcessor	Async gross-to-net in 500-employee chunks; saga compensation on retry
finance-outbox	ApService	OutboxProcessor	Durable domain event delivery (Outbox pattern)
notification-dispatch	NotificationService	NotificationDispatchProcessor	Email/SMS/webhook — 5 retries, exponential back-off, DLQ in Bull Board
forecast-retrain	Cron scheduler	ForecastRetrainProcessor	Weekly ML model re-training + Redis cache invalidation

6 Security & Compliance

6.1 Authentication & Token Lifecycle

- RS256 JWT from Keycloak, verified per request against the realm JWKS endpoint.
- **Token blacklist:** logout writes the JTI to Redis; every guarded route checks the set — guarantees token invalidation even within the access-token TTL.
- **Realm-per-tenant:** tenant slugs are immutable; a cross-realm token cannot authenticate to another tenant’s API because the JWKS URL differs.
- **MFA:** per-tenant TOTP enforcement via Keycloak’s `Configure OTP required action` — triggers on both password and SSO login paths.

6.2 Role-Based Access Control

78 @Roles() declarations across 21 controllers. The matrix below is abridged; the full matrix is in docs/rbac-role-matrix.md.

Operation	SA	TA	Mgr	Vwr	Emp
Finance (GL / AP / AR) — read	✓	✓	✓	✓	—
Finance — write / approve	✓	✓	✓	—	—
Fiscal period close	✓	✓	—	—	—
HR — employee CRUD	✓	✓	✓	—	—
Payroll run	✓	✓	✓	—	—
Leave request	✓	✓	✓	✓	✓
Leave approve/reject	✓	✓	✓	—	—
SCM — PO / receive	✓	✓	✓	—	—
Employee delete	✓	✓	—	—	—
Audit log / GDPR	✓	✓	—	—	—

SA = SuperAdmin, TA = TenantAdmin, Mgr = Manager, Vwr = Viewer, Emp = Employee

6.3 Transport & Headers

Caddy terminates TLS (Let’s Encrypt, auto-renewed) for all public services. The API sets: Strict-Transport-Security (max-age 1 year), X-Frame-Options: DENY, X-Content-Type-Options: nosniff, Content-Security-Policy (nonce-based script-src), and Permissions-Policy. CORS is locked to the exact frontend origin.

6.4 Rate Limiting & Throttling

@nestjs/throttler is configured with user-aware keys (logged-in user id rather than raw IP to survive NAT). A k6 load test at 2,000 VUs surfaced a JWKS caching bug (re-fetching on every request) and a throttler misconfiguration that rejected legitimate traffic — both fixed before the final load test run.

6.5 Static Analysis & Supply-Chain Scans

The CI pipeline (.github/workflows/ci.yml) includes:

- **Tenant-scoping gate:** a custom TypeScript-AST script that fails the build if any Prisma query is missing tenantId in its where.
- **TruffleHog:** secret scanning on every push.
- **Grype + Trivy:** vulnerability scanning of Docker images and node_modules.
- **ESLint security plugin:** runs as a lint-staged hook — no warnings can be committed.

6.6 Key Vulnerabilities Found & Fixed

Finding	Severity	Fix
Dead tenant middleware with header fallback	High	Deleted unused <code>tenant-context.middleware.ts</code> ; confirmed the live interceptor uses only the verified JWT
IDOR via raw <code>findUnique</code> without <code>tenantId</code>	High	Full grep audit + CI gate; all 38 services audited, compound unique constraints added where needed
Duplicate AP invoice from dual event dispatch	High	Removed BullMQ path; <code>goods.received</code> processed by exactly one in-process listener
JWKS re-fetch per request	Medium	Caching with TTL added to <code>KeycloakStrategy</code>
Throttler keyed on IP (NAT bypass)	Medium	Switched to user-aware key using authenticated user ID

7 Quality Assurance & Performance

7.1 API Testing — 64/64 Authenticated Checks

A Postman/Newman suite of 64 end-to-end API scenarios runs against the live deployment for each of the six business flows (auth/tenant, GL, AP, HR, SCM, PM). Every test authenticates with a real Keycloak token; failure aborts the CI pipeline. As of 13 July 2026 all 64 pass.

7.2 Unit Tests — 31 CI-Enforced Tests

31 unit tests written with **Vitest** cover the deterministic money paths, run on every push via the `unit-tests` CI job:

- `invoice-matching.service.spec.ts` — **17 scenarios**: 10 covering the header-total fallback (exact match, tolerance boundary inclusive at 2%, symmetric under/overcharge, vendor mismatch, wrong GR, missing PO/GR, zero-total PO never auto-approves) and 7 covering the newer line-level path (partial-receipt invoicing, quantity-exceeds-received, per-line price tolerance, unresolved product, one bad line fails the whole invoice).
- `payroll-deductions.spec.ts` — **9 scenarios**: gross-to-net under and at the ESI ceiling, overtime included in gross/TDS base but PF on basic only, gratuity as an employer-only accrual, never-negative net pay, and four tax-slab resolution cases (default fallback, range match, open-ended top slab, inclusive boundaries).
- `gl.service.spec.ts` — **5 scenarios**: balanced entry posts, unbalanced entry rejected and never written, multi-line aggregate balance accepted, posting into a locked fiscal period rejected, posting into a non-existent period rejected.

The payroll saga's crash/retry compensation is verified separately, live against the real BullMQ worker and database rather than as an isolated unit test — see `pnpm -filter api verify:payroll-retry` and Section 10.

7.3 Load Testing — k6 at 2,000 VUs

The first three runs used a single shared login and immediately exposed two real bugs: a JWKS client re-created per request (saturating Keycloak) and a rate limiter keyed by raw IP (all 2,000 virtual users collapsed into one bucket behind NAT). Both looked identical in symptom — near-instant request failure, not slow timeouts — which is what pointed at rejection rather than genuine overload. After fixing both, Run 4 used a pool of **150 distinct provisioned users** so the fix could be demonstrated under real concurrent load rather than a single shared token:

Metric	Run 1 (baseline, broken)	Run 4 (fixed, 150-user pool)
Checks succeeded	10.04%	98.01%
http_req_duration P95	39 ms (<i>on the 10% that got a real response</i>)	1.48 s (<i>genuine queueing under real load, not rejection</i>)

The honest reading: fixing the two bugs took success from 10% to 98% at true 2,000-VU concurrency. The 1.48s P95 does not meet the TDD's 300ms target at this scale — load testing surfaced a genuine, unfixed finding: **a single Node.js process is a real throughput ceiling at 2,000 concurrent users**, requiring clustering or horizontal scaling (an infrastructure change, not a code fix). This is recorded as a roadmap item rather than glossed over.

7.4 Lighthouse Audits

Page	Performance	Accessibility	Best Practices	SEO
/login	90	96	100	100
/home	94	100	96	100
/bi	82	90	96	100

/bi performance is lower due to eager chart-library loading. Before fixes, /home and /bi started at 63 and 40 respectively — raised by a silent-SSO redirect fix and dynamic imports on 3 pages (48% First Load JS cut).

7.5 Database Query Optimisation

Every table had only a handful of seed rows in development, so slow queries were invisible by construction. To get a real signal, **80,000 synthetic rows** were seeded (20k AuditLog, 20k Notification, plus Invoice, JournalEntry, PurchaseOrder, StockMovement) and the actual query each service issues was run through EXPLAIN ANALYZE before and after adding composite indexes:

Table	Before	After (index added)
AuditLog	6.26 ms — bitmap heap scan + sort	0.22 ms — index scan, sort eliminated
Invoice	2.56 ms — sequential scan + sort	0.28 ms — <i>needed a 3-column index (below)</i>
Notification	4.99 ms — bitmap heap scan + sort	0.19 ms — index scan, sort eliminated
JournalEntry	2.56 ms — bitmap heap scan + sort	0.17 ms — index scan, sort eliminated
StockMovement	3.04 ms — bitmap AND (2 indexes) + sort	0.42 ms — single composite index scan

Indexes added: (tenantId, createdAt) on JournalEntry, PurchaseOrder, and PurchaseRequisition; (tenantId, productId, warehouseId) on InventoryCostLayer; (tenantId, productId, createdAt) on StockMovement. The Invoice case caught a real planning mistake: a 2-column (tenantId, type) index was *not used by the planner at all* — with type = 'AP' matching roughly half the seeded rows, Postgres correctly preferred a sequential scan over an index scan that still needed a separate sort for ORDER BY createdAt. Adding createdAt as a third column ((tenantId, type, createdAt)) let the index satisfy the filter *and* the sort together — confirmed only because the fix was verified with EXPLAIN ANALYZE rather than applied speculatively. Separately, an N+1 query in ap.service.ts's createPaymentRun() (a per-invoice findFirst loop) was replaced with one upfront findMany plus an in-memory lookup map.

8 Deployment & DevOps

8.1 Production Environment

The platform is deployed on a single **Oracle Cloud Ampere VM** (4 OCPU, 24 GB RAM, Ubuntu 22.04). The production stack:

Component	Details
Reverse proxy	Caddy 2 — automatic Let’s Encrypt TLS, zero-downtime reload
Process manager	pm2 — <code>amdox-api</code> (port 3101) + <code>amdox-web</code> (port 3000)
Data services	Docker Compose — PostgreSQL 17, Redis 8, Elasticsearch 9, MinIO, Keycloak 25
Observability	Docker Compose (separate project) — OTel Collector, Tempo, Prometheus, Loki, Promtail, Grafana
Public URLs	<code>erp.</code> , <code>kc.</code> , <code>grafana.</code> , <code>mailpit.92-4-86-3.sslip.io</code> — app, Keycloak, dashboards, alert inbox

8.2 CI Pipeline

The GitHub Actions pipeline (`.github/workflows/ci.yml`) runs 10 jobs on every push:

1. **Lint** — ESLint across the monorepo, zero warnings tolerated.
2. **Typecheck** — `tsc -noEmit` for both apps.
3. **Unit tests (Vitest)** — the 31-test suite; fails the build on any regression.
4. **Tenant-scoping audit** — the custom AST script; fails on any `tenantId`-less Prisma query or raw `PrismaClient` usage.
5. **Build** — `nest build` + Next.js production build.
6. **Secret scan** — TruffleHog across all staged files.
7. **Dependency scan** — Gripe, with SARIF upload to GitHub Security.
8. **Container scan** — Trivy filesystem scan (dependencies + misconfig) across the whole repo.
9. **Image scan** — Trivy image scan per service (`api/web/ml-service` matrix).
10. **Slack notification** — posts the pipeline result to the team channel.

A separate post-deployment smoke test (health/readiness checks for `api`, `web`, and `ml-service`) runs as an ArgoCD PostSync hook after a Kubernetes rollout — not as a GitHub Actions job.

8.3 Container & Kubernetes

All three services (`api`, `web`, `ml-service`) are built as **multi-stage distroless Docker images** to minimise attack surface. A **Helm chart** defines the Kubernetes deployment: `HorizontalPodAutoscaler` (CPU >70%), `PodDisruptionBudget`, Ingress with TLS, and an **Istio canary** setup (stable/canary `DestinationRule` subsets + a weighted `VirtualService`) validated with 1,000 real in-mesh requests splitting 889/111 against a 90/10 target. **ArgoCD** watches a Git source and auto-syncs the cluster with self-heal enabled — demonstrated by a config change deployed via a plain `git commit`, a manual `kubect1 edit` self-healing back within ~20 seconds, and a post-install/ PostSync hook Job that runs health/readiness smoke checks (`api`, `web`, `ml-service`) before a rollout is considered complete. A local `kind` cluster hosts three namespaces — `amdox-dev`, `amdox-staging`, `amdox-prod` — each from the same chart with per-env values files, distinct `JWT_SECRETS`, and host-based ingress routing; Sealed Secrets encrypt credentials at rest (verified: a restarted pod passed readiness against a controller-decrypted `DATABASE_URL`).

9 Observability

The full observability stack was implemented and verified live on 12 July 2026:

Signal	Implementation
Traces	OTel SDK in <code>apps/api/src/observability/otel.ts</code> ; HTTP, Prisma, Redis, BullMQ auto-instrumented. Tail sampling: 100% errors, 10% successes. Stored in Tempo; browsable in Grafana Explore.
Metrics	Prometheus scrapes <code>:9464/metrics</code> every 15 s. The “Amdox API — Golden Signals” dashboard: request rate, 5xx %, p50/p95 latency (300 ms SLA line), host CPU/RAM, event-loop utilisation, live log panel.
Logs	Promtail ships pm2 logs (<code>job="pm2"</code>) and Docker logs (<code>job="docker"</code>) to Loki. Searchable in Grafana Explore with LogQL.
Alerts	Two SLA rules: 5xx > 5% for 5 min; p95 > 300 ms for 5 min. Both email to Mailpit (verified with a test notification). PagerDuty/OpsGenie is a receiver swap.

Public access: grafana.92-4-86-3.sslip.io (admin / `Amdox-Obs-2026!`) and mailpit.92-4-86-3.sslip.io (amdox / `Amdox-Obs-2026!`). Prometheus, Loki, and Tempo remain localhost-only (no built-in auth); their data surfaces through Grafana’s authenticated UI.

One deployment gotcha: Oracle Cloud’s default iptables rejects container-to-host traffic. Opening ports 9464 and 1025 to the Docker subnet only, and persisting the rules in `/etc/iptables/rules.v4`, resolved the metrics and SMTP connectivity issues.

10 Engineering Challenges & Notable Fixes

Challenge	Root Cause	Fix
Duplicate AP invoices & doubled GL postings	<code>receiveGoods</code> dispatched <code>goods.received</code> to <i>both</i> a BullMQ queue and an in-process <code>EventEmitter</code> . Both consumers raced past the idempotency check and each created an approved invoice.	Removed the BullMQ path entirely; one in-process listener now owns invoice creation. Data cleanup script executed on existing records.
Payroll double-payslip on BullMQ retry	A failed run's partial payslips were never deleted before the automatic retry. Restart from chunk 0 re-inserted all payslips already written, inflating <code>totalNetPay</code> and the GL entry.	Processor now deletes the failed run's payslips before reprocessing (saga compensation). Proven by a kill-and-retry test: 21 employees → exactly 21 payslips before and after.
JWKS re-fetch on every API request	<code>KeycloakStrategy.validate()</code> called the JWKS endpoint without caching; at 2,000 VUs this became the bottleneck and saturated the Keycloak container.	Added in-process JWKS cache with a 5-minute TTL using <code>jwt-rsa</code> caching options.
Throttler rejecting legitimate traffic at load	Default <code>@nestjs/throttler</code> keyed on raw IP. Behind NAT, all 2,000 VUs shared one IP and exhausted the limit instantly.	Switched to user-aware throttler key (authenticated user ID) post-JWT validation.
ES SDK v9 incompatible with ES 8.x server	The <code>npm</code> package <code>@elastic/elasticsearch@9</code> sends <code>compatible-with=9</code> headers; ES 8.x rejects them.	Upgraded the Docker container from ES 8.15 to ES 9.0.3 (fresh volume; all data in Postgres, ES is a cache).
Keycloak Google login blocked by CONFIG-URE_TOTP	A test user in the <code>company-z</code> realm had <code>CONFIGURE_TOTP</code> as a required action. Keycloak's IDP post-login flow cannot prompt for required actions and failed with <code>IDENTITY_PROVIDER_POST_LOGIN_ERROR</code> .	Removed the required action from the user via the Keycloak Admin API. Per-tenant MFA toggle now manages OTP enrollment cleanly.
Elasticsearch query DSL strict TypeScript types	ES SDK v9 uses a discriminated union for <code>QueryDslQueryContainer</code> where the <code>bool.must</code> array type is inferred as <code>never</code> in certain overloads.	Wrapped <code>client.search</code> in a private <code>searchIndex</code> helper with <code>(this.client.search as any)</code> to bypass the strict SDK types while preserving full runtime behaviour.
Oracle iptables blocking container→host metrics	Oracle's default firewall rules reject traffic from Docker bridge networks to host ports (9464/metrics, 1025/Mailpit).	Opened specific ports to the Docker subnets only; rules persisted in <code>/etc/iptables/rules.v4</code> .

Table 2: Selected engineering challenges encountered and resolved during the project.

11 Requirements Compliance & Known Deviations

11.1 Functional Requirements Scoreboard

Req	Name	Status	Notes / Deviations
F-01	Multi-tenant auth (SSO)	Done	Realm-per-tenant Keycloak; OIDC + Google auto-link; per-tenant MFA toggle; SAML tested
F-02	Financial ledger (GL)	Done	Double-entry + period-close enforced; auto-posting; ECB FX rates fetched but not yet applied in postings (roadmap)
F-03	AP/AR automation	Done	Textract OCR (mock fallback); line-level 3-way match; partial receipts; payment runs; aging report
F-04	HR & payroll	Done	Chunked BullMQ engine; PF/ESI/PT/TDS; payslip PDFs; retry bug fixed (verify:payroll-retry); leave accrual cron
F-05	SCM & inventory	Done	Requisition→PO→GR (partial receipts); FIFO cost layers; reorder automation; vendor portal
F-06	AI demand forecasting	Done	Prophet + PyTorch LSTM; Redis cache; weekly retrain; MAPE breach alerting
F-07	Project management	Done	DAG tasks; Gantt; utilisation heatmap; budget actuals from AP + payroll; overrun alerts
F-08	Business intelligence	Done	Drag-drop grid; 10+ chart types; SSE live refresh; scheduled PDF/XLSX reports
F-09	Audit & compliance	Done	SHA-256 hash-chained audit log + verify; GDPR DSR + consent; Postgres (TimescaleDB deviation)
F-10	Notification engine	Done	In-app SSE + email + SMS + HMAC webhook; per-user preferences; 5-retry queue; Bull Board DLQ
F-11	API gateway & webhooks	Done	REST api/v1 + Swagger; GraphQL (Apollo); Elasticsearch 10 entity types; outbound HMAC webhooks
F-12	Offline / PWA	Open	De-scoped for MVP per TDD Day-1 guidance; on roadmap

11.2 Non-Functional Requirements — Evidence

NFR	Target	Evidence
API latency	<300 ms P95	k6 (150-user pool, 2,000 VU): 98.01% success, P95 = 1.48 s — honest finding: a single-process throughput ceiling, not yet met at this scale (Section 7)
Throughput	2,000 concurrent	k6 2,000-VU/10-min run; JWKS caching + IP-keyed throttler bugs found and fixed; single-Node scaling ceiling found, roadmap
Security	OWASP Top 10	Helmet/CSP, user-aware rate limiting, CI tenant-scoping gate, Truffle-Hog/Grype/Trivy
Availability	99.9%	Single-VM + Caddy; SLA alert rules in Grafana; HA = roadmap
Frontend	Lighthouse ≥90	/login 90, /home 94, /bi 82 (performance; eager chart loading is the remaining /bi lever)

11.3 Technology Stack — Spec vs. Delivered

Category	TDD Specifies	Delivered	Match
Frontend	Next.js 15 + React 19 + TS	Next.js 15.5 + React 19.1	✓
UI library	shadcn/ui + Tailwind 4	Custom component kit + Tailwind 4	≈
Backend	NestJS 11	NestJS 11 modular monolith	✓
Database	PostgreSQL 17 + Prisma	PostgreSQL 17 + Prisma (67 models)	✓
Cache/Queue	Redis 8 + BullMQ	Redis 8 + BullMQ (4 queues)	✓
ML	FastAPI + Prophet + LSTM	FastAPI + Prophet + PyTorch LSTM	✓
Search	Elasticsearch 8	Elasticsearch 9 (upgraded)	✓
Auth	Keycloak 25 OIDC/SAML	Keycloak 25, realm-per-tenant, MFA, SAML tested	✓
Containers	Docker multi-stage distroless	All 3 services, distroless, verified	✓
Orchestration	Kubernetes + Helm	Helm chart + ArgoCD on kind (local cluster)	✓
Observability	OTel + Prometheus + Grafana	Full stack + Tempo + tail sampling + SLA alerts	✓
IaC	Terraform	Not implemented (single-VM, Helm per env)	×
Time-series	TimescaleDB	Plain Postgres + hash chain for audit	≈

11.4 Known Deviations

1. **F-12 PWA/offline:** the only un-implemented functional requirement. Service worker and cache-first strategy are on the roadmap.
2. **Multi-currency conversion:** ECB rates are fetched and stored daily but are not yet applied in journal postings or invoice totals. The data model and FX service are in place; conversion logic is the remaining piece.
3. **Terraform IaC:** the current deployment is a single VM managed by pm2 and Docker Compose. Terraform modules for Oracle Cloud would be added if expanding to a multi-region or multi-node setup.
4. **shadcn/ui:** the TDD recommended shadcn/ui; a custom Tailwind component kit was used instead, achieving the same visual result with tighter control over bundle size.

12 Learnings & Conclusion

12.1 Key Technical Learnings

Working on a production ERP from scratch over a structured 28-day plan required reasoning about the full software lifecycle — from schema design and API contracts through to CI pipelines, observability stacks, and live incident resolution. The most impactful individual learnings were:

- **Domain events as a modularity primitive:** routing cross-module work through EventEmitter2 rather than direct service calls enforced segregation of duties at the code level. Finance genuinely cannot be called from SCM — the compiler prevents it.
- **Idempotency is not optional in event-driven systems:** the duplicate AP-invoice bug (two consumers racing past the same idempotency check) made clear that “fire-and-forget” events need guard rails at the consumer side, not just the emitter.
- **Observability must be built in, not bolted on:** the Prometheus + Tempo + Loki stack revealed the JWKS caching bottleneck during load testing that would have been invisible from application logs alone.
- **Multi-tenancy has multiple enforcement layers:** a single `tenantId` column is necessary but not sufficient. The combination of JWT-scoped identity, a CI static-analysis gate, and the Prisma

- \$extends wrapper is what makes tenant isolation provably correct rather than merely hopeful.
- **Statutory payroll is precise:** implementing PF, ESI, PT, LWF, and income-tax TDS correctly required reading the actual Indian Labour Law regulations rather than approximate percentages. The statutory compliance config table allows per-tenant override for different states or future rule changes.

12.2 Roadmap

Item	Description
F-12 PWA / Offline	Service worker with cache-first API responses and IndexedDB sync queue
Multi-currency posting	Apply stored ECB rates in GL postings, invoice totals, and report currency conversion
Terraform IaC	Oracle Cloud VCN, compute, and load-balancer resources as code
BI bundle splitting	Code-split remaining eagerly-loaded chart libraries to raise the /bi Lighthouse performance score further
Recursive-CTE org chart	Server-side tree traversal for large employee hierarchies (currently assembled client-side from <code>managerId</code>)
Horizontal API scaling	Address the single-Node throughput ceiling found under k6 load (clustering / multi-instance)
shadcn/ui migration	Adopt the TDD-specified component library for full design-system parity

12.3 Conclusion

The **Amdox AI-Powered Cloud ERP Suite** delivers a working, publicly deployed, multi-tenant enterprise application that meets 11 of the 12 functional requirements set by the company’s technical design document. The codebase is production-grade in the sense that matters most: the money paths (GL double-entry, 3-way match, payroll gross-to-net, saga compensation) are unit-tested and load-tested; security findings discovered during the audit were fixed and proven in code; and the observability stack provides SLA alerting and distributed tracing that a real operations team would rely on.

This project demonstrated that a well-structured internship can produce deliverables that go meaningfully beyond tutorial-level code: real Keycloak realm provisioning, real FIFO inventory costing, real BullMQ saga patterns, and real k6 bugs caught in production. The remaining gaps (PWA, FX conversion, Terraform) are honest deviations from the original plan — documented, understood, and tracked — not hidden defects.

Pawan Nannaware
Industry Intern — Amdox Technologies
July 2026 • github.com/PawanYN/amdox-erp